

ZAYNOR Forensic Report — NITROBA

Overview — what kind of incident this is

Classification: SUSPICION — no ATT&CK; technique corroborated

Confidence: UNKNOWN

Result SHA-256: 434cf0a51eaa0649e7a78a8941dfe4d3d6c9ce16c99b403c706e03889cb9089b

Engine: ZAYNOR deterministic engine (Mode 1) 1.0.0-SANS-2026

Findings / Unknowns: 1 / 0

Agents in this pipeline

Role	Status	What it actually does
MENTOR	conectado	The only role a normal zaynor chat/serve run actually invokes; explains an already-sealed result, never re-invokes the engine.
INVESTIGATOR	implementado, sin invocación automática	collect_window/verify_custody are real and tested (agents/investigator_tools.py) but no CLI/API command calls them yet.
FLEET_COMMANDER	implementado, sin invocación automática	Writes to the investigation log; real code (agents/fleet_commander_tools.py), no caller in cli.py/api.py yet.
DETECTION_ENGINEER	implementado, sin invocación automática	draft_sigma_rule anchors candidates to a real sealed finding (agents/detection_engineer_tools.py), not wired to a command yet.
DISPATCHER	implementado, sin invocación automática	Catalog of evidence types Mode 1 can actually analyze (agents/dispatcher_tools.py), not wired to a command yet.
ENDPOINT_HUNTER / PERSISTENCE_HUNTER	out of scope	Would need a live EDR collection backend this project does not have.
THREAT_INTEL	out of scope, for now	A portable VirusTotal/GTI enrichment exists but is not wired in — external network dependency, pending decision.

Findings come from ZAYNOR's own deterministic Mode 1 engine, not from a named agent above -- no per-finding agent attribution exists in the sealed contract.

Findings

Finding	State	MITRE	Detected by	Rationale
F-NITROBA	SUSPICION	-	ZAYNOR deterministic engine (Mode 1)	[MOTOR] Label-blind selection: SUSPICION_DETECTED. Significant signal with structural support (score=0.3135)

Unknowns

- None declared.

Chain of custody

```
result_sha256 (deterministic):
434cf0a51eaa0649e7a78a8941dfe4d3d6c9ce16c99b403c706e03889cb9089b
report_hash (timestamped):
000dcbd7157a9cb373e9a67c54f8b14af2ea666b7c0f527892c4055373a3ab0b
```

Methodology

The deterministic engine produces and seals the result before any language model is invoked. The model receives a compressed, read-only summary and narrates it; it cannot alter a verdict, a finding, or the seal. Every claim in a narration is checked against this sealed result before being shown (zaynor chat/serve). `result_sha256` is the canonical SHA-256 seal of the authoritative result: recompute it independently with ``zaynor audit`` to confirm this report was not altered after the fact.